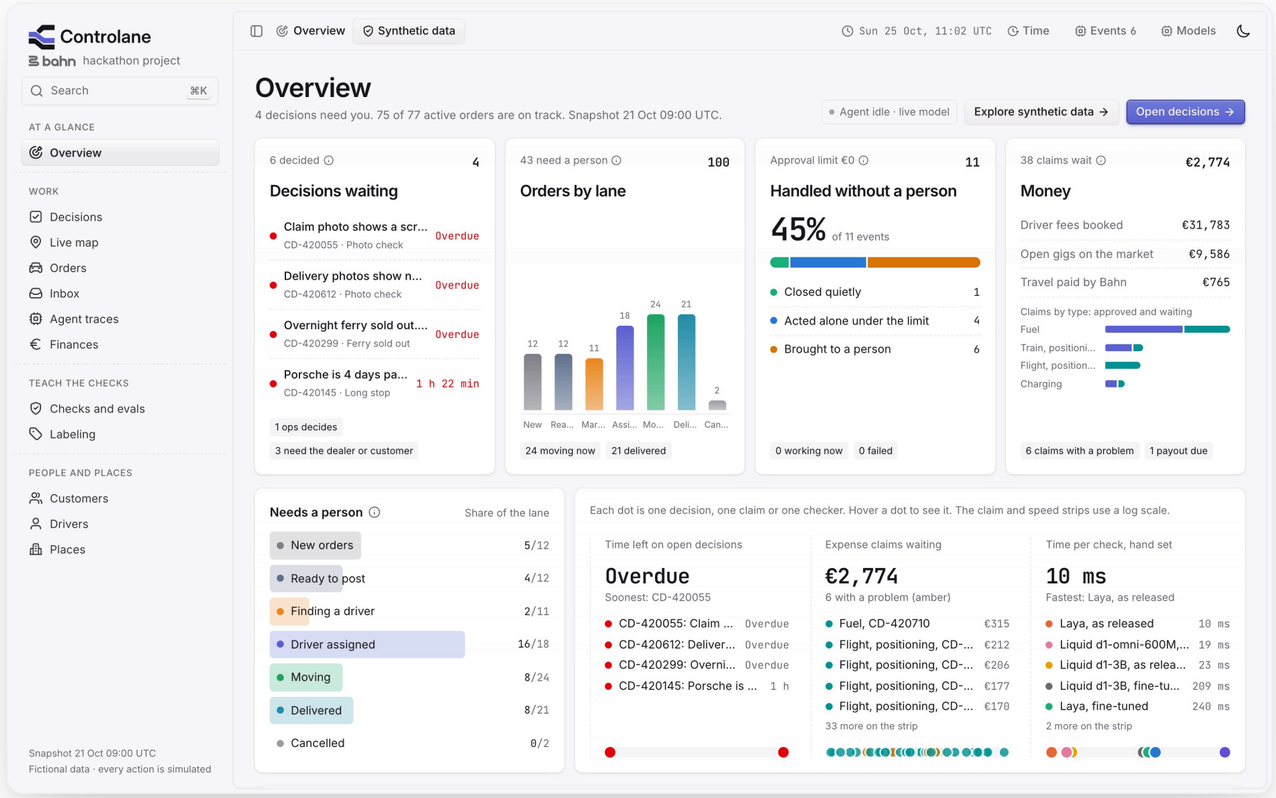
Controlane

Agents do the legwork, people decide

Veikka Silvekoski · bob the builder

controlane.bobs.build



The problem

Every change starts a chain someone has to hold

Each problem is simple alone. The chain is the hard part.

25,000+
cars this year

1,500+
drivers

about 50
people in operations

3
countries

Step 1
Ferry sold out

→

Step 2
New booking

→

Step 3
Hotel?

→

Step 4
Email to the ferry line

→

Step 5
Cost decision

→



Today, one person holds all of it

Source: Bahn challenge brief

The control room

75 of 77 active orders need nobody; 4 decisions come first

100 orders

77 active

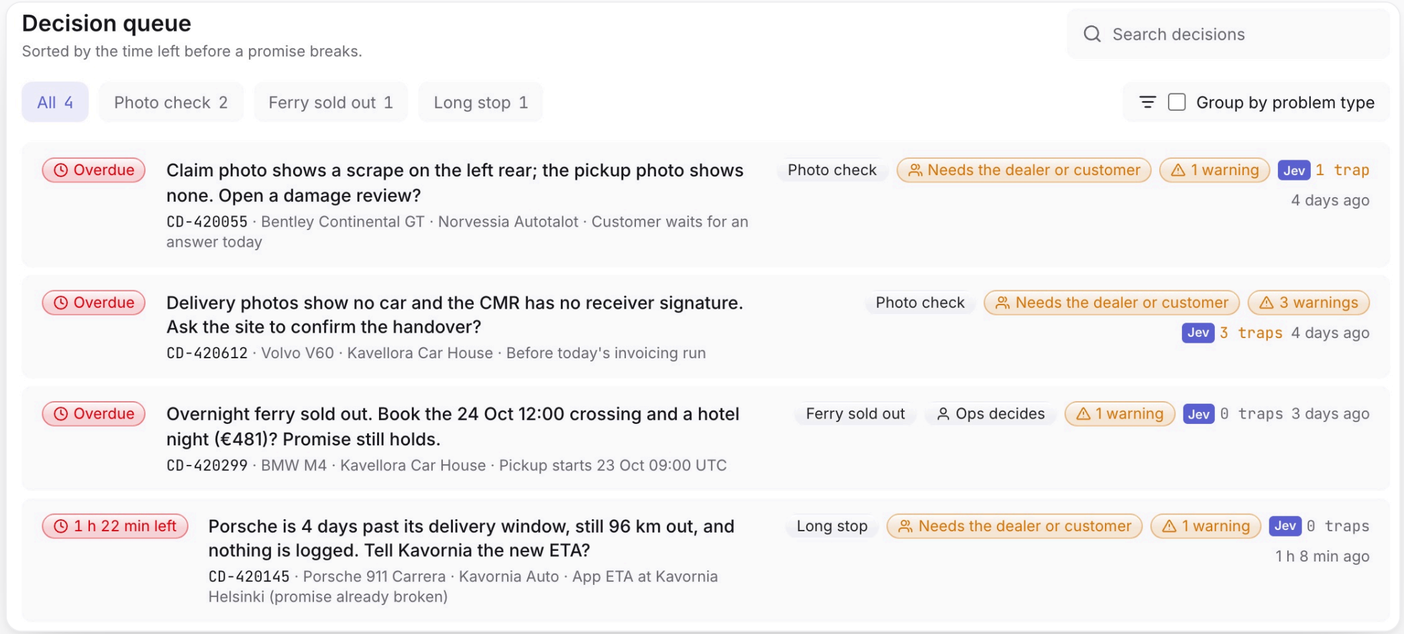
23 delivered or cancelled leave here

75 on track

the agents handle these

4 decisions

a person decides these first

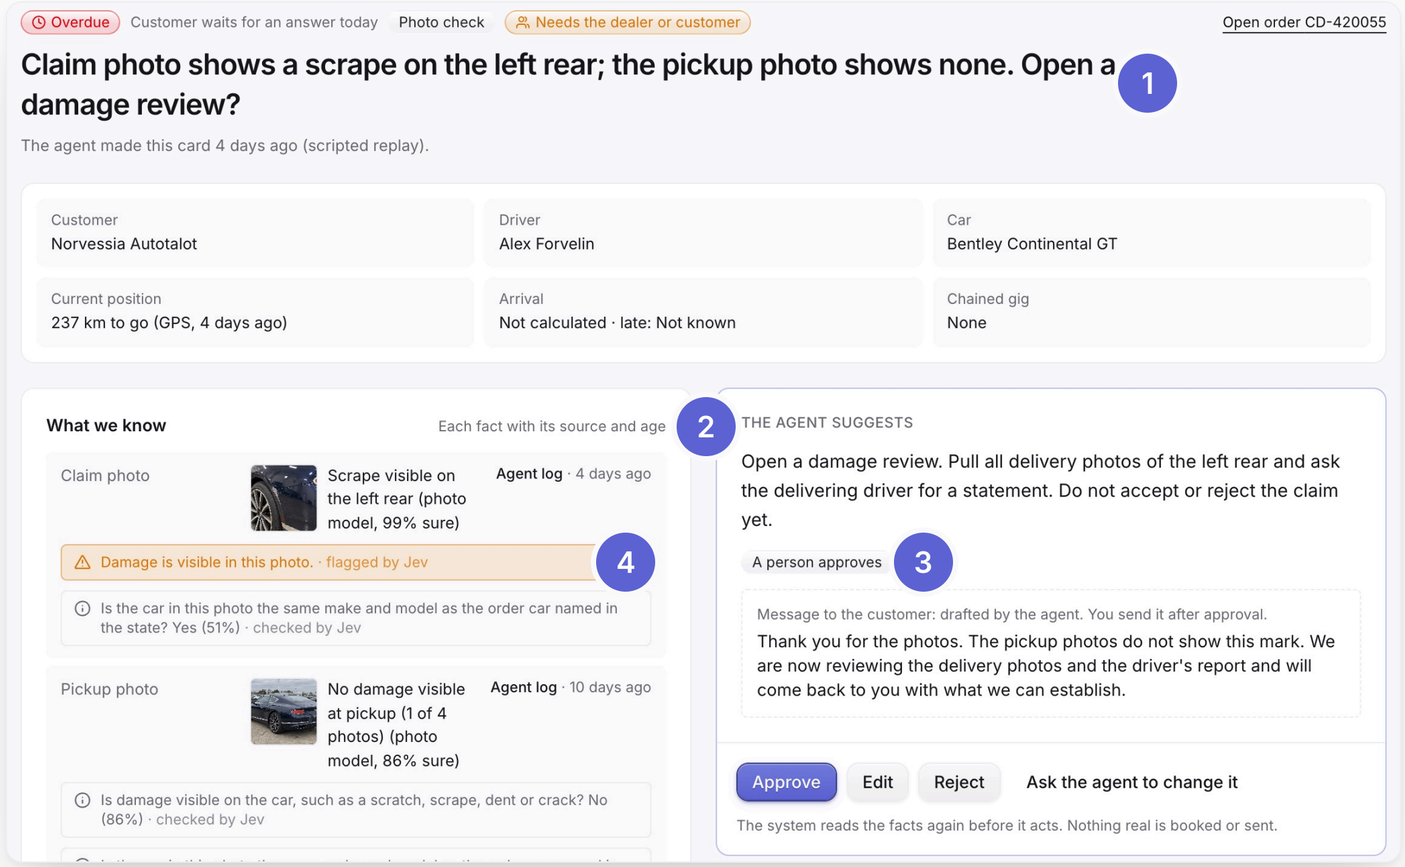


“Sorted by the time left before a promise breaks.” From the app

Source: overview snapshot, 21 Oct 09:00 UTC · Fictional data, every action is simulated

• The control room

One card shows what happened, who said it, and what is next



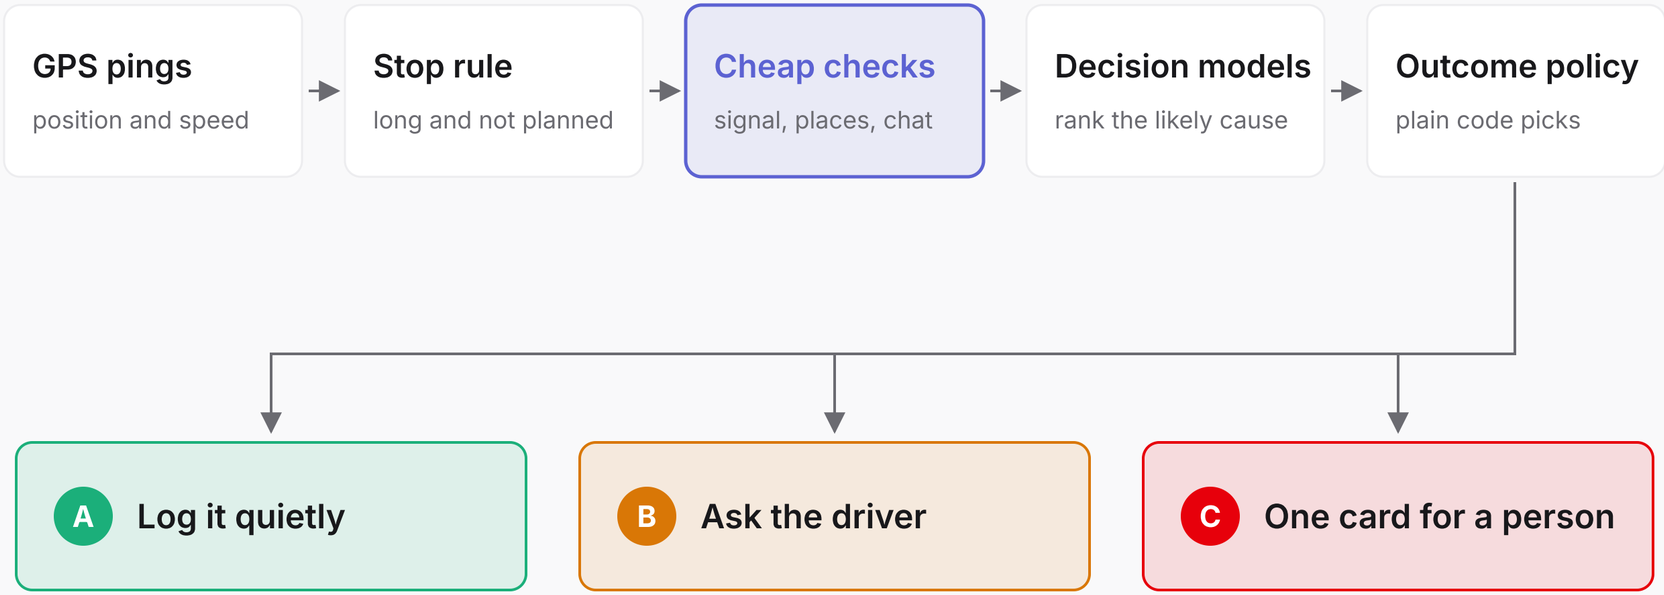
1. Question first, and who decides
2. Every fact with its source and age
3. The suggestion, and who approves it
4. Models flag known traps before you approve

“The system reads the facts again before it acts.” From the app · Fictional data, every action is simulated

The agent

The stop check does the cheap work first

A plain rule finds the stop. The models rank the cause. Plain code picks the outcome.

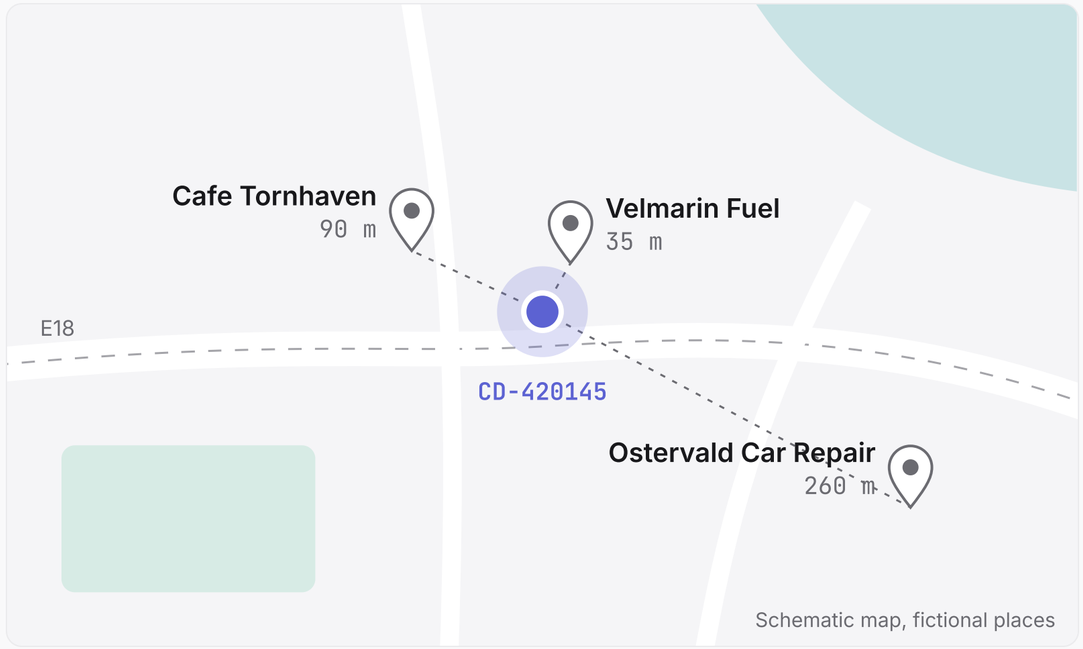


CHEAP FIRST

GPS signal	free
Places lookup	one call
Driver chat	free
Arrival time	free
Decision models	about 26 ms

The control room

One stop, three outcomes, one decision for a person



- **09:47** Car stopped
- **10:08** A: logged quietly
- **10:22** B: asked the driver
- **10:27** Reply: warning light, and a photo
- **10:29** C: one card for a person

THE CARD SUGGESTS “Call the dealer: their own roadside help, or should we arrange a tow?”

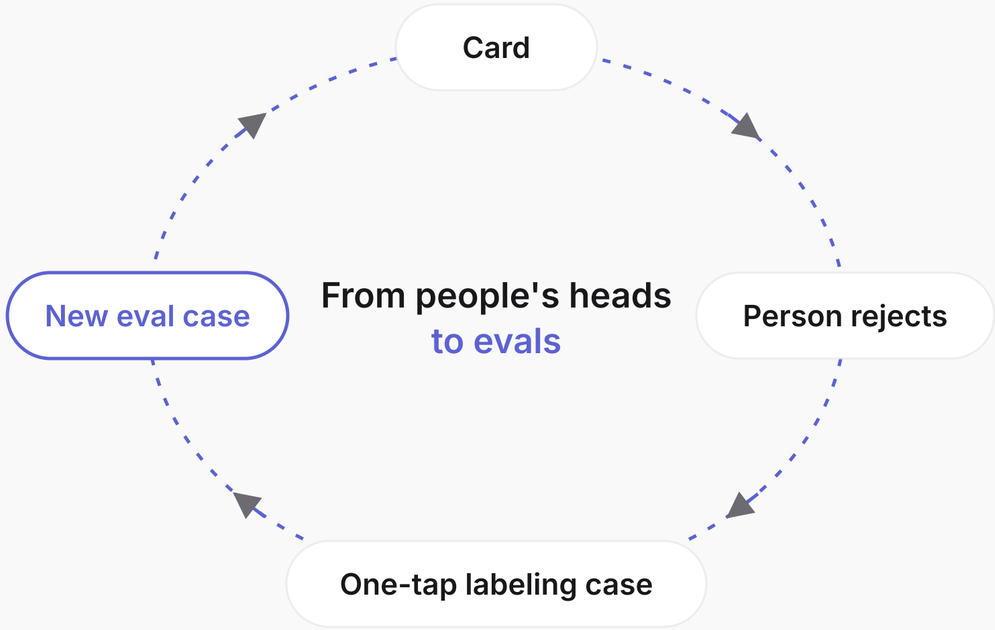
Source: the worked example in the stop check build spec · Fictional data, every action is simulated

The control room

A person approves every change, and every rejection becomes an eval

Look up and read	The agent alone
Message a driver	The agent alone
Message a customer	A person approves
Book, spend or move a date	A person approves

More freedom only when a person raises the limit.



The model

A $1.33 fine-tune beats the commercial checker on unseen cases

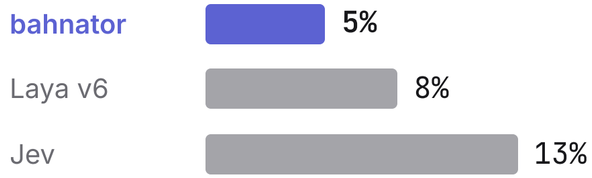
bahnator-1.1-pro-preview

Liquid d1-3B + LoRA rank 16 · 24.5 M trainable parameters · 7 training runs, 2 models

Correct answers



False alarms, lower is better



Milliseconds per check, lower is better



$1.33 · 24 min training · 1x RTX 6000 Ada · Verda, Finland

[bahnator weights on Hugging Face] [Laya v6 weights on Hugging Face]

Synthetic test: 4,395 cases from 203 template families held out of training

How I built it

PRD first, four models, one cross-AI review

PRD.md

Safe On-the-Road Agent

Problem

Rules and approvals

Interface

Benchmark against Jev

Visual checks


Overview


Decision card


Live map


Checks and evals


Labeling

My Claude framework on Opus 5.5

[skill name] [skill name] [skill name]

Pi coding agent

Gemini 3.8 Flash | Grok 4.7 | GPT-6 Astra

→ **Cross-AI review**

BUILT WITH

WHERE I DID NOT TAKE THE FIRST ANSWER

Whole template families held out, not random rows, so the test measures new kinds of cases.

[Add one more true example]

Try it

Try it at controlane.bobs.build



controlane.bobs.build

GitHub
[GitHub repo URL]

bahnator weights
[bahnator weights on Hugging Face]

Laya v6 weights
[Laya v6 weights on Hugging Face]

NEXT

1. A real places lookup behind the mock
2. More eval families from rejected cards
3. More freedom per action as approval rates prove out